

# # Lab 4: Real-Time Data Streaming with Amazon Kinesis & CloudWatch

## Step 1: Create a Kinesis Data Stream

**Data stream configuration**

**Data stream name**

realtime-sales-stream

Acceptable characters are uppercase and lowercase letters, numbers, underscores, hyphens and periods.

**Data stream capacity** [Info](#)

**Capacity mode**

☐ On-demand  
Use this mode when your data stream's throughput requirements are unpredictable and variable. With on-demand mode, your data stream's capacity scales automatically.

☒ Provisioned  
Use provisioned mode when you can reliably estimate throughput requirements of your data stream. With provisioned mode, your data stream's capacity is fixed.

**Provisioned shards**

The total capacity of a stream is the sum of the capacities of its shards. Enter number of provisioned shards to see total data stream capacity.

1 [Shard estimator](#)

Minimum: 1, Maximum available: 500, Account quota limit: 500. [Request shard quota increase](#)

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

## Step 2: Send Data to Kinesis Stream (AWS CLI)

```
~ $ aws kinesis put-record \  
> --stream-name realtime-sales-stream \  
> --partition-key "partition1" \  
> --data "$(echo -n '{"sale_id": 101, "amount": 250.75, "timestamp": "'$(date -u +%Y-%m-%dT%H:%M:%SZ)'" }' | base64)" \  
> --region us-east-1  
{  
  "ShardId": "shardId-000000000000",  
  "SequenceNumber": "49661391800661484331231243076634560561890376342143238146"  
}
```

```
~ $ aws kinesis put-record --stream-name realtime-sales-stream --partition-key "partition1" --data "$(echo -n '{"sale_id": 101, "amount": 250.75, "timestamp": "'$(date -u +%Y-%m-%dT%H:%M:%SZ)'" }' | base64)" --region us-east-1  
{  
  "ShardId": "shardId-000000000000",  
  "SequenceNumber": "49661391800661484331231243076634560561890376342143238146"  
}  
~ $ aws kinesis put-record --stream-name realtime-sales-stream --partition-key "partition1" --data "$(echo -n '{"sale_id": 101, "amount": 250.75, "timestamp": "'$(date -u +%Y-%m-%dT%H:%M:%SZ)'" }' | base64)" --region us-east-1  
{  
  "ShardId": "shardId-000000000000",  
  "SequenceNumber": "49661391800661484331231243557411060838612675241665626114"  
}  
~ $ aws kinesis put-record --stream-name realtime-sales-stream --partition-key "partition1" --data "$(echo -n '{"sale_id": 101, "amount": 250.75, "timestamp": "'$(date -u +%Y-%m-%dT%H:%M:%SZ)'" }' | base64)" --region us-east-1  
{  
  "ShardId": "shardId-000000000000",  
  "SequenceNumber": "4966139180066148433123124356289474835638463372788654338"  
}  
~ $ aws kinesis put-record --stream-name realtime-sales-stream --partition-key "partition1" --data "$(echo -n '{"sale_id": 101, "amount": 250.75, "timestamp": "'$(date -u +%Y-%m-%dT%H:%M:%SZ)'" }' | base64)" --region us-east-1  
{  
  "ShardId": "shardId-000000000000",  
  "SequenceNumber": "49661391800661484331231243575716615599217391716895031298"  
}
```

### Step 3: Create a Lambda Function to Process Stream Data

The screenshot displays the AWS Lambda console interface. At the top, the navigation bar shows the user is in the 'United States (N. Virginia)' region. The main content area is titled 'Create function' and offers three options: 'Author from scratch' (selected), 'Use a blueprint', and 'Container image'. Below this, the 'Basic information' section is visible, showing the function name 'kinesis-sales-processor', the runtime 'Python 3.13', and the architecture 'x86\_64'. A green notification banner at the top of the code editor states 'Successfully updated the function kinesis-sales-processor.' The code editor shows a Python file named 'lambda\_function.py' with the following code:

```
1 import json
2
3 def lambda_handler(event, context):
4     for record in event["Records"]:
5         # Decode Kinesis data (base64)
6         payload = json.loads(record["kinesis"]["data"].decode('utf-8'))
7         print(f"Processed sale: {payload}")
8     return {"statusCode": 200, "body": "Success"}
```

The interface also includes a sidebar with 'EXPLORER' and 'DEPLOY' sections, and a bottom status bar with 'CloudShell' and 'Feedback' links.

### Step 4:

Attached Lambda to Kinesis Stream and in trigger I Select realtime-sales-stream in Lambda settings

## Step 5: Monitor Kinesis Stream Data in CloudWatch Logs Insights

The screenshot shows the AWS CloudWatch console interface. The left sidebar contains navigation links for CloudWatch, Favorites and recents, Dashboards, AI Operations, Alarms, Logs, Metrics, and X-Ray traces. The main content area displays the configuration for a log group named `/aws/lambda/kinesis-sales-processor`. The configuration includes fields for Creation time (6 minutes ago), Retention (Never expire), Stored bytes, Contributor insights, KMS key ID, Anomaly detection, and Transform. Below the configuration, there are tabs for Log streams, Tags, Anomaly detection, Metric filters, Subscription filters, Contributor Insights, and Data protection. The Log streams tab is active, showing a list of log streams with columns for Log stream and Last event time. A search bar is available to filter log streams.

The screenshot shows the AWS CloudWatch Logs Insights console interface. The left sidebar contains navigation links for CloudWatch, Favorites and recents, Dashboards, AI Operations, Alarms, Logs, Metrics, and X-Ray traces. The main content area displays a query editor with a query: `fields @timestamp, @message | sort @timestamp desc | limit 20`. Below the query editor, there are buttons for Run query, Cancel, Save, and History. A status message indicates that the query is completed and executed for 1 log group. Below the status message, there are tabs for Logs (10), Patterns (5), and Visualization. The Logs tab is active, showing a list of logs with columns for Log stream and Last event time. A search bar is available to filter logs. A histogram is also displayed, showing the distribution of logs over time.

| #   | @timestamp               | @message   |
|-----|--------------------------|--|
| ▶ 1 | 2025-03-14T01:58:11.555Z | REPORT RequestId: 574094b7-6de5-4652-b110-b6790d5fa503 Duration: 1.68 ms Billed Duration: 2 ms Memory Size: 128 MB Max Mem |
| ▶ 2 | 2025-03-14T01:58:11.555Z | END RequestId: 574094b7-6de5-4652-b110-b6790d5fa503  |
| ▶ 3 | 2025-03-14T01:58:11.541Z | Processed sale: {'sale_id': 101, 'amount': 250.75, 'timestamp': '2025-03-14T00:38:53Z'}                                    |
| ▶ 4 | 2025-03-14T01:58:11.540Z | START RequestId: 574094b7-6de5-4652-b110-b6790d5fa503 Version: \$LATEST  |
| ▶ 5 | 2025-03-14T01:58:11.531Z | REPORT RequestId: 910985dd-bc82-4938-b6c9-219116a443ab Duration: 2.29 ms Billed Duration: 3 ms Memory Size: 128 MB Max Mem |
| ▶ 6 | 2025-03-14T01:58:11.530Z | END RequestId: 910985dd-bc82-4938-b6c9-219116a443ab  |
| ▶ 7 | 2025-03-14T01:58:11.529Z | Processed sale: {'sale_id': 101, 'amount': 250.75, 'timestamp': '2025-03-14T00:38:49Z'}                                    |
| ▶ 8 | 2025-03-14T01:58:11.528Z | START RequestId: 910985dd-bc82-4938-b6c9-219116a443ab Version: \$LATEST  |

## Step 6: Cleanup

### 1. Delete Kinesis Stream:

✔ realtime-sales-stream was successfully deleted.

Data streams (1) [Info](#)

Process data in real time

Create a Firehose stream

Actions ▾

Create data stream

Find data streams

< 1 >

⚙

| <input type="checkbox"/> | Name                                  | Status  | Capacity mode | Provisioned shards | Sharing policy | Data retention period | Encryption | Consumers with enhanced fan-out |
|--------------------------|---------------------------------------|---------|---------------|--------------------|----------------|-----------------------|------------|---------------------------------|
| <input type="checkbox"/> | <a href="#">realtime-sales-stream</a> | Deleted | -             | -                  | -              | -                     | -          | -                               |

### 2. Delete Lambda Function:

Functions (5)

Last fetched 0 seconds ago

Actions ▾

Create function

Filter by attributes or search by keyword

< 1 > ⚙

| <input type="checkbox"/> | Function name                             | Description   | Package type | Runtime    | Last modified |
|--------------------------|---|---|--------------|------------|---------------|
| <input type="checkbox"/> | <a href="#">RoleCreationFunction</a>      | Create SLR if absent                                  | Zip          | Python 3.9 | 3 days ago    |
| <input type="checkbox"/> | <a href="#">MainMonitoringFunction</a>    | -   | Zip          | Python 3.9 | 3 days ago    |
| <input type="checkbox"/> | <a href="#">RedshiftEventSubscription</a> | Create Redshift event subscription to SNS Topic.      | Zip          | Python 3.9 | 3 days ago    |
| <input type="checkbox"/> | <a href="#">ModLabRole</a>                | updates LabRole to allow it to assume itself          | Zip          | Python 3.9 | 3 days ago    |
| <input type="checkbox"/> | <a href="#">RedshiftOverwatch</a>         | Deletes Redshift Cluster if the count is more than 2. | Zip          | Python 3.9 | 3 days ago    |

### 3. Delete CloudWatch Log Group:

aws

Search

[Alt+S]

United States (N. Virginia)

voclabs/user3839096=rbarboz@stevens.edu @ 6901-7535-0910

CloudWatch > Log groups

CloudWatch

Favorites and recents

Dashboards

✔ The following log group(s) have been deleted:

- /aws/lambda/kinesis-sales-processor

Filter log groups or try prefix search

Exact match

< 1 > ⚙