

1) I created an EC2 instance using the Amazon Linux 2 AMI to serve as my web server. This instance was selected for its compatibility with the Nginx web server and its efficiency in handling network traffic.

▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - *required*

Royce_lab5 ▼

↻ Create new key pair

▼ Network settings [Info](#)

Network [Info](#)

vpc-04b23a6a95db17cc0

Subnet [Info](#)

No preference (Default subnet in any availability zone)

Auto-assign public IP [Info](#)

Enable

Additional charges apply when outside of free tier allowance

Firewall (security groups) [Info](#)

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

○ Create security group

○ Select existing security group

We'll create a new security group called 'launch-wizard-3' with the following rules:

☒ Allow SSH traffic from

Helps you connect to your instance

Anywhere
0.0.0.0/0 ▼

☐ Allow HTTPS traffic from the internet

To set up an endpoint, for example when creating a web server

☒ Allow HTTP traffic from the internet

To set up an endpoint, for example when creating a web server

Edit

▼ Summary

Number of instances [Info](#)

5

When launching more than 1 instance, [consider EC2 Auto Scaling](#)

Software Image (AMI)

Amazon Linux 2 Kernel 5.10 AMI...[read more](#)

ami-0a9a48ce4458e384e

Virtual server type (instance type)

t2.micro

Firewall (security group)

New security group

Storage (volumes)

1 volume(s) - 8 GiB

❗ Free tier: In your first year of opening an AWS account, you get 750 hours per month of t2.micro instance usage (or t3.micro where t2.micro isn't available) when used with free tier AMIs, 750 hours per month of public IPv4 address usage, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of bandwidth to the internet.

Cancel

Launch instance

[Preview code](#)

[Alt+S] ⓘ 🔔 ⌚ ⚙️ United States (N. Virginia) ▼ voclabs/user3839096=rbarboz@stevens.edu @ 6901-7535-0910

Instances (1/5) [Info](#)

Last updated less than a minute ago ↻

Connect

Instance state ▼

Actions ▼

Launch instances

▼

🔍 Find Instance by attribute or tag (case-sensitive)

All states ▼

< 1 > ⚙️

<input type="checkbox"/>	Name ↗ ▼	Instance ID	Instance state ▼	Instance type ▼	Status check	Alarm status	Availability Zone
<input type="checkbox"/>	Load Balancer	i-031bfcd56196e942d	✔ Running ⓘ ⓘ	t2.micro	⌚ Initializing	View alarms +	us-east-1c
<input type="checkbox"/>	Server1	i-0b41737a90e23d1af	✔ Running ⓘ ⓘ	t2.micro	⌚ Initializing	View alarms +	us-east-1c
<input type="checkbox"/>	Server2	i-0b338c069982e7926	✔ Running ⓘ ⓘ	t2.micro	⌚ Initializing	View alarms +	us-east-1c
<input type="checkbox"/>	Server3	i-0d1c5d69ccb934c74	✔ Running ⓘ ⓘ	t2.micro	⌚ Initializing	View alarms +	us-east-1c
<input checked="" type="checkbox"/>	Server4	i-0b87449644e3155c8	✔ Running ⓘ ⓘ	t2.micro	⌚ Initializing	View alarms +	us-east-1c

◀ ▶

2) connected to Server 1,2,3,4 using SSH, and I installed Nginx using the command `sudo yum install nginx -y`. This setup will allow the server to handle incoming HTTP requests from the load balancer.

```
Royce Barboz@royce-asus MINGW64 ~/OneDrive/Desktop/cloud (main)
$ chmod 400 Royce_lab5.pem

Royce Barboz@royce-asus MINGW64 ~/OneDrive/Desktop/cloud (main)
$ ssh -i "Royce_lab5.pem" ec2-user@ec2-54-85-54-122.compute-1.amazonaws.com
Last login: Sun Apr  6 18:07:39 2025 from pool-100-8-48-95.nwrknj.fios.verizon.net

#_
~\  #####_      Amazon Linux 2
~~\  #####\
~~\  #####|
~~\  #/
~~\  V~'-'>
~~~
~~~_.._
~~~_/_/_/_/_/_
    _/m/'-/_/_/_/_/_

AL2 End of Life is 2026-06-30.

A newer version of Amazon Linux is available!

Amazon Linux 2023, GA and supported until 2028-03-15.
https://aws.amazon.com/linux/amazon-linux-2023/

[ec2-user@ip-172-31-91-181 ~]$ sudo amazon-linux-extras install nginx1 -y
Installing nginx
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
Cleaning repos: amzn2-core amzn2extra-docker amzn2extra-kernel-5.10
                  : amzn2extra-nginx1
17 metadata files removed
6 sqlite files removed
0 metadata files removed
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
amzn2-core                                     | 3.6 kB  00:00:00
amzn2extra-docker                             | 2.9 kB  00:00:00
amzn2extra-kernel-5.10                       | 3.0 kB  00:00:00
amzn2extra-nginx1                           | 2.9 kB  00:00:00
(1/9): amzn2-core/2/x86_64/group_gz         | 2.7 kB  00:00:00
(2/9): amzn2-core/2/x86_64/updateinfo       | 1.0 MB  00:00:00
(3/9): amzn2extra-docker/2/x86_64/updateinfo | 23 kB  00:00:00
(4/9): amzn2extra-docker/2/x86_64/primary_db | 124 kB  00:00:00
(5/9): amzn2extra-nginx1/2/x86_64/updateinfo | 3.9 kB  00:00:00
(6/9): amzn2extra-nginx1/2/x86_64/primary_db | 61 kB  00:00:00
(7/9): amzn2extra-kernel-5.10/2/x86_64/updateinfo | 121 kB  00:00:00
(8/9): amzn2extra-kernel-5.10/2/x86_64/primary_db | 35 MB  00:00:01
(9/9): amzn2-core/2/x86_64/primary_db       | 75 MB  00:00:01
Resolving Dependencies
--> Running transaction check
--> Package nginx.x86_64 1:1.26.3-1.amzn2.0.1 will be installed
--> Processing Dependency: nginx-core = 1:1.26.3-1.amzn2.0.1 for package: 1:nginx-1.26.3-1.amzn2.0.1.x86_64
--> Processing Dependency: nginx-filesystem = 1:1.26.3-1.amzn2.0.1 for package: 1:nginx-1.26.3-1.amzn2.0.1.x86_64

[ec2-user@ip-172-31-91-181 ~]$ sudo systemctl start nginx
[ec2-user@ip-172-31-91-181 ~]$ sudo systemctl enable nginx
Created symlink from /etc/systemd/system/multi-user.target.wants/nginx.service to /usr/lib/systemd/system/nginx.service.
```

Not secure http://ec2-54-85-54-122.compute-1.amazonaws.com

https://github.com/... AWS Skill Builder Learner Dashboard... 30 Days Coding GitHub - monajalal/..

Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.


```
[ec2-user@ip-172-31-91-181 ~]$ cd /usr/share/nginx/html  
[ec2-user@ip-172-31-91-181 html]$ sudo nano index.html
```

GNU nano 2.9.8


```
<!DOCTYPE html>  
<html>  
<head>  
<title>Welcome to nginx!</title>  
<style>  
html { color-scheme: light dark; }  
body { width: 35em; margin: 0 auto;  
font-family: Tahoma, Verdana, Arial, sans-serif; }  
</style>  
</head>  
<body>  
<h1>Welcome to nginx!</h1>  
<p>If you see this page, the nginx web server is successfully installed and  
working. Further configuration is required.</p>  
  
<p>For online documentation and support please refer to  
<a href="http://nginx.org/">nginx.org</a>.<br/>  
Commercial support is available at  
<a href="http://nginx.com/">nginx.com</a>.</p>  
  
<p><em>Thank you for using nginx.</em></p>  
</body>  
</html>
```

```
ec2-user@ip-172-31-91-181:/usr/share/nginx/html
GNU nano 2.9.8


<h1>Server1</h1>
```




Server1



Server2




Server3



Server4

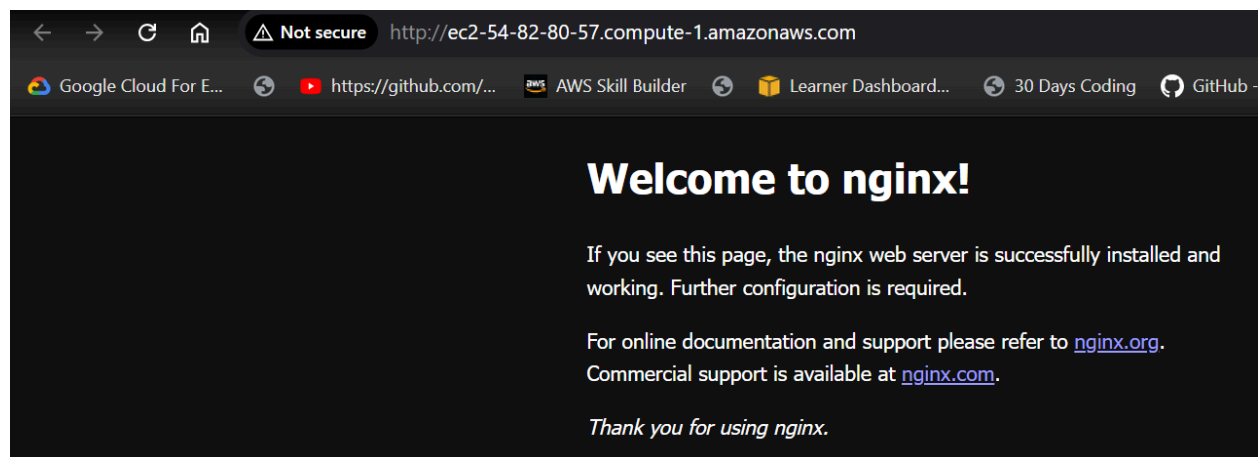
3) I configured the Nginx load balancer to distribute traffic across multiple servers. This was done by editing the nginx.conf file and defining an upstream block with multiple servers and weighted load balancing. Each server was assigned a specific weight to control the traffic distribution

```
Royce Barboz@royce-asus MINGW64 ~/OneDrive/Desktop/cloud (main)
$ ssh -i "Royce_lab5.pem" ec2-user@ec2-54-82-80-57.compute-1.amazonaws.com
The authenticity of host 'ec2-54-82-80-57.compute-1.amazonaws.com (54.82.80.57)' can't be established.
ED25519 key fingerprint is SHA256:RyEWB9aDHnySqjmOIQkqKQT5d1kkCRV/qwejs0IzvK.
This key is not known by any other names
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-54-82-80-57.compute-1.amazonaws.com' (ED25519) to the list of known hosts.
```



```
#_
##### Amazon Linux 2
#####
####| AL2 End of Life is 2026-06-30.
#/#
V~'|_>
~~~~
A newer version of Amazon Linux is available!
~~~~
.._|/_/___/___/
_/m/'      Amazon Linux 2023, GA and supported until 2028-03-15.
           https://aws.amazon.com/linux/amazon-linux-2023/
```

```
[ec2-user@ip-172-31-88-116 ~]$ sudo amazon-linux-extras install nginx1 -y
Installing nginx
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
Cleaning repos: amzn2-core amzn2extra-docker amzn2extra-kernel-5.10
```



```
70 unbound1.17          available [ =stable ]
72 collectd-python3     available [ =stable ]
† Note on end-of-support. Use 'info' subcommand.
[ec2-user@ip-172-31-88-116 ~]$ sudo systemctl start nginx
[ec2-user@ip-172-31-88-116 ~]$ sudo systemctl enable nginx
Created symlink from /etc/systemd/system/multi-user.target.wants/nginx.service to /usr/lib/systemd/system/nginx.service.
```

```
ec2-user@ip-172-31-88-116:~
GNU nano 2.9.8

# For more information on configuration, see:
# * Official English Documentation: http://nginx.org/en/docs/
# * Official Russian Documentation: http://nginx.org/ru/docs/

user nginx;
worker_processes auto;
error_log /var/log/nginx/error.log;
pid /run/nginx.pid;

# Load dynamic modules. See /usr/share/doc/nginx/README.dynamic.
include /usr/share/nginx/modules/*.conf;

events {
    worker_connections 1024;
}

http {
    log_format main '$remote_addr - $remote_user [$time_local] "$request" '
        '$status $body_bytes_sent "$http_referer" '
        '"$http_user_agent" "$http_x_forwarded_for"';

    access_log /var/log/nginx/access.log main;

    sendfile            on;
    tcp_nopush          on;
    tcp_nodelay         on;
    keepalive_timeout   65;
    types_hash_max_size 4096;

    include              /etc/nginx/mime.types;
    default_type         application/octet-stream;

    # Load modular configuration files from the /etc/nginx/conf.d directory.
    # See http://nginx.org/en/docs/nginx\_core\_module.html#include
    # for more information.
    include /etc/nginx/conf.d/*.conf;

    server {
        listen          80;
        listen          [::]:80;
        server_name     _;
        root             /usr/share/nginx/html;
    }
}
```

```
ec2-user@ip-172-31-88-116:~  
GNU nano 2.9.8  
  
events {  
    worker_connections 768;  
}  
  
http {  
    upstream myapp {  
        server ec2-54-85-54-122.compute-1.amazonaws.com weight=1;  
        server ec2-13-219-86-50.compute-1.amazonaws.com weight=2;  
        server ec2-54-161-55-70.compute-1.amazonaws.com weight=3;  
        server ec2-3-95-19-202.compute-1.amazonaws.com weight=4;  
    }  
  
    server {  
        listen 80;  
        server_name myapp.com;  
  
        location / {  
            proxy_pass http://myapp;  
        }  
    }  
}
```

4. After configuring the load balancer, I tested its functionality by accessing the load balancer's public DNS. The responses were correctly routed to different web servers, verifying that the load balancer was distributing traffic as intended.

```
[ec2-user@ip-172-31-88-116 ~]$ sudo systemctl restart nginx  
[ec2-user@ip-172-31-88-116 ~]$ curl ec2-54-82-80-57.compute-1.amazonaws.com  
<h1>Server4</h1>  
[ec2-user@ip-172-31-88-116 ~]$ curl ec2-54-82-80-57.compute-1.amazonaws.com  
<h1>Server3</h1>  
[ec2-user@ip-172-31-88-116 ~]$ curl ec2-54-82-80-57.compute-1.amazonaws.com  
<h1>Server2</h1>  
[ec2-user@ip-172-31-88-116 ~]$ curl ec2-54-82-80-57.compute-1.amazonaws.com  
<h1>Server4</h1>  
[ec2-user@ip-172-31-88-116 ~]$ |
```

5. I created a Python script to simulate 2000 visits to the load balancer. This allowed me to collect data on how the load balancer distributed traffic across the web servers. The script successfully recorded the number of visits to each server.

```
ec2-user@ip-172-31-88-116:~  
GNU nano 2.9.8  
  
import requests  
from collections import Counter  
  
# Define the load balancer's DNS name (replace with your load balancer's DNS)  
load_balancer_dns = "http://ec2-54-82-80-57.compute-1.amazonaws.com"  
  
# Create a Counter to store the visit counts  
visit_counts = Counter()  
  
# Simulate 2000 visits to the load balancer  
for _ in range(2000):  
    try:  
        # Send a request to the load balancer  
        response = requests.get(load_balancer_dns)  
  
        # Check which server was accessed based on the response text  
        if 'Server1' in response.text:  
            visit_counts['Server1'] += 1  
        elif 'Server2' in response.text:  
            visit_counts['Server2'] += 1  
        elif 'Server3' in response.text:  
            visit_counts['Server3'] += 1  
        elif 'Server4' in response.text:  
            visit_counts['Server4'] += 1  
    except requests.RequestException as e:  
        print(f"Error with request: {e}")  
        continue  
  
# Display the summary of the visit counts  
print("Visit Counts:")  
for server, count in visit_counts.items():  
    print(f"{server} visit counts: {count}")  
print(f"Total visit counts: {sum(visit_counts.values())}")
```

Scenario 1: Weighted Traffic Distribution (Increasing Weights for Servers)

In this scenario, the weights of the servers were set as follows: Server 1 (weight 1), Server 2 (weight 2), Server 3 (weight 3), and Server 4 (weight 4). As a result, the traffic distribution was not equal. Server 4, with the highest weight, handled the most traffic, followed by Server 3, Server 2, and Server 1. The load balancer routed the majority of requests to Server 4, which handled approximately 40% of the total traffic, with the other servers handling traffic in proportion to their assigned weights. This confirmed that the load balancer was accurately respecting the weight distribution.


```
ec2-user@ip-172-31-88-116:~  
GNU nano 2.9.8  
  
events {  
    worker_connections 768;  
}  
  
http {  
    upstream myapp {  
        server ec2-54-85-54-122.compute-1.amazonaws.com weight=1;  
        server ec2-13-219-86-50.compute-1.amazonaws.com weight=2;  
        server ec2-54-161-55-70.compute-1.amazonaws.com weight=3;  
        server ec2-3-95-19-202.compute-1.amazonaws.com weight=4;  
    }  
  
    server {  
        listen 80;  
        server_name myapp.com;  
  
        location / {  
            proxy_pass http://myapp;  
        }  
    }  
}
```

```
[ec2-user@ip-172-31-88-116 ~]$ python3 visit_collector.py  
Visit Counts:  
Server4 visit counts: 800  
Server3 visit counts: 600  
Server2 visit counts: 400  
Server1 visit counts: 200  
Total visit counts: 2000  
[ec2-user@ip-172-31-88-116 ~]$
```

Scenario 2: Equal Traffic Distribution (Weight = 1 for All Servers)

In this scenario, each server received an equal amount of traffic, as all servers were assigned the same weight (1). The results showed that the load balancer distributed the incoming traffic evenly across all four servers. Each server handled approximately 25% of the total traffic, demonstrating that the load balancer was working as expected to equally balance the load.

ec2-user@ip-172-31-88-116:~

GNU nano 2.9.8

```
events {
    worker_connections 768;
}

http {
    upstream myapp {
        server ec2-54-85-54-122.compute-1.amazonaws.com weight=1;
        server ec2-13-219-86-50.compute-1.amazonaws.com weight=1;
        server ec2-54-161-55-70.compute-1.amazonaws.com weight=1;
        server ec2-3-95-19-202.compute-1.amazonaws.com weight=1;
    }

    server {
        listen 80;
        server_name myapp.com;

        location / {
            proxy_pass http://myapp;
        }
    }
}
```

[ec2-user@ip-172-31-88-116 ~]\$ sudo systemctl restart nginx

[ec2-user@ip-172-31-88-116 ~]\$ python3 visit_collector.py

Visit Counts:

Server1 visit counts: 500

Server2 visit counts: 500

Server3 visit counts: 500

Server4 visit counts: 500

Total visit counts: 2000

[ec2-user@ip-172-31-88-116 ~]\$ |

Scenario 3: Mixed Weight Distribution (Weight = 1, 2, 1, 2)

For this scenario, the weights were set as follows: Server 1 (weight 1), Server 2 (weight 2), Server 3 (weight 1), and Server 4 (weight 2). The traffic distribution was not as even as Scenario 1, but the traffic was distributed in a manner that reflected the weights. Servers 2 and 4, with higher weights, received a larger share of the traffic, while Servers 1 and 3 handled fewer requests. This scenario demonstrated the load balancer's ability to allocate traffic based on weighted values, with Servers 2 and 4 handling a larger portion of the load compared to Servers 1 and 3.

```
ec2-user@ip-172-31-88-116:~  
GNU nano 2.9.8  
  
events {  
    worker_connections 768;  
}  
  
http {  
    upstream myapp {  
        server ec2-54-85-54-122.compute-1.amazonaws.com weight=1;  
        server ec2-13-219-86-50.compute-1.amazonaws.com weight=2;  
        server ec2-54-161-55-70.compute-1.amazonaws.com weight=1;  
        server ec2-3-95-19-202.compute-1.amazonaws.com weight=2;  
    }  
  
    server {  
        listen 80;  
        server_name myapp.com;  
  
        location / {  
            proxy_pass http://myapp;  
        }  
    }  
}
```

```
[ec2-user@ip-172-31-88-116 ~]$ sudo nano /etc/nginx/nginx.conf  
[ec2-user@ip-172-31-88-116 ~]$ sudo systemctl restart nginx  
[ec2-user@ip-172-31-88-116 ~]$ python3 visit_collector.py  
Visit Counts:  
Server2 visit counts: 667  
Server4 visit counts: 667  
Server1 visit counts: 333  
Server3 visit counts: 333  
Total visit counts: 2000  
[ec2-user@ip-172-31-88-116 ~]$
```

6. Used the `script` command to record all the commands I used

```
GNU nano 2.9.8                                     ec2_creation_log.txt

Script started on 2025-04-06 21:01:37+0000
^[[0;ec2-user@ip-172-31-88-116:~$ aws configure
AWS Access Key ID [None]: 690175350910
AWS Secret Access Key [None]: royce
Default region name [None]: us-east-1
Default output format [None]: json
^[[0;ec2-user@ip-172-31-88-116:~$ aws ec2 run-instances --image-id ami-0a9a48ce4458e384e --count 1 --instance-type t3.micro
An error occurred (AuthFailure) when calling the RunInstances operation: AWS was not able to validate the provided access credentials
^[[0;ec2-user@ip-172-31-88-116:~$ sudo tcpdump -r traffic.pcap
tcpdump: traffic.pcap: No such file or directory
^[[0;ec2-user@ip-172-31-88-116:~$ sudo yum install tcpdump -y
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
MAMzn2-core
Package 14:tcpdump-4.9.2-4.amzn2.1.0.i.x86_64 already installed and latest version
Nothing to do
^[[0;ec2-user@ip-172-31-88-116:~$ sudo tcpdump -i eth0 -w traffic.pcap
tcpdump: listening on eth0, link-type EN10MB (Ethernet), capture size 262144 bytes
^C86 packets captured
86 packets received by filter
0 packets dropped by kernel
^[[0;ec2-user@ip-172-31-88-116:~$ sudo tcpdump -i eth0 -w traffic.pcap
reading from file traffic.pcap, link-type EN10MB (Ethernet)
21:18:29.609014 IP ip-172-31-88-116.ec2.internal.ssh > pool-100-35-52-54.nwrknj.fios.verizon.net.62694: Flags [P.], seq 1745134480:1745134524, win 452, len 0
21:18:29.609066 IP ip-172-31-88-116.ec2.internal.ssh > pool-100-35-52-54.nwrknj.fios.verizon.net.62694: Flags [P.], seq 44:160, ack 1, win 452, len 0
21:18:29.627129 IP pool-100-35-52-54.nwrknj.fios.verizon.net.62694 > ip-172-31-88-116.ec2.internal.ssh: Flags [.], ack 160, win 253, length 0
21:18:32.840595 ARP, Request who-has ip-172-31-88-116.ec2.internal tell ip-172-31-88-116.ec2.internal, length 28
21:18:32.840609 ARP, Reply ip-172-31-88-116.ec2.internal is-at 12:10:b5:7c:97:e9 (oui Unknown), length 28
21:18:39.143594 IP ip-172-31-88-116.ec2.internal.50356 > rn-02.koehn.com.ntp: NTPv4, Client, length 48
21:18:39.159883 IP rn-02.koehn.com.ntp > ip-172-31-88-116.ec2.internal.50356: NTPv4, Server, length 48
21:18:40.786224 IP ip-172-31-88-116.ec2.internal.42120 > 169.254.169.123.ntp: NTPv4, Client, length 48
21:18:40.786925 IP 169.254.169.123.ntp > ip-172-31-88-116.ec2.internal.42120: NTPv4, Server, length 48
21:18:56.839582 IP ip-172-31-88-116.ec2.internal.53121 > 169.254.169.123.ntp: NTPv4, Client, length 48
21:18:56.840127 IP 169.254.169.123.ntp > ip-172-31-88-116.ec2.internal.53121: NTPv4, Server, length 48
21:19:02.047522 ARP, Request who-has ip-172-31-80-1.ec2.internal tell ip-172-31-88-116.ec2.internal, length 28
21:19:02.047820 ARP, Reply ip-172-31-80-1.ec2.internal is-at 12:74:3c:d2:04:55 (oui Unknown), length 28
21:19:11.137569 IP pool-100-35-52-54.nwrknj.fios.verizon.net.64009 > ip-172-31-88-116.ec2.internal.http: Flags [S], seq 976386369, win 65535, length 0
21:19:11.137593 IP ip-172-31-88-116.ec2.internal.http > pool-100-35-52-54.nwrknj.fios.verizon.net.64009: Flags [S.], seq 1544814859, ack 976386369, win 65535, length 0
21:19:11.145862 IP pool-100-35-52-54.nwrknj.fios.verizon.net.64010 > ip-172-31-88-116.ec2.internal.http: Flags [S], seq 2708787082, win 65535, length 0
21:19:11.145872 IP ip-172-31-88-116.ec2.internal.http > pool-100-35-52-54.nwrknj.fios.verizon.net.64010: Flags [S.], seq 3242955271, ack 2708787082, win 65535, length 0
21:19:11.153600 IP pool-100-35-52-54.nwrknj.fios.verizon.net.64009 > ip-172-31-88-116.ec2.internal.http: Flags [.], ack 1, win 255, length 0
21:19:11.163135 IP pool-100-35-52-54.nwrknj.fios.verizon.net.64010 > ip-172-31-88-116.ec2.internal.http: Flags [.], ack 1, win 255, length 0
21:19:13.030471 IP ip-172-31-88-116.ec2.internal.33732 > 169.254.169.123.ntp: NTPv4, Client, length 48
```

AG Get Help	AR Write Out	AW Where Is	AK Cut Text	AJ Justify	AC Cur Pos	M-U Undo	M-A Mark Text	M-W What
AX Exit	AR Read File	AA Replace	AK Uncut Text	AT To Spell	AG Go To Line	M-E Redo	M-G Copy Text	


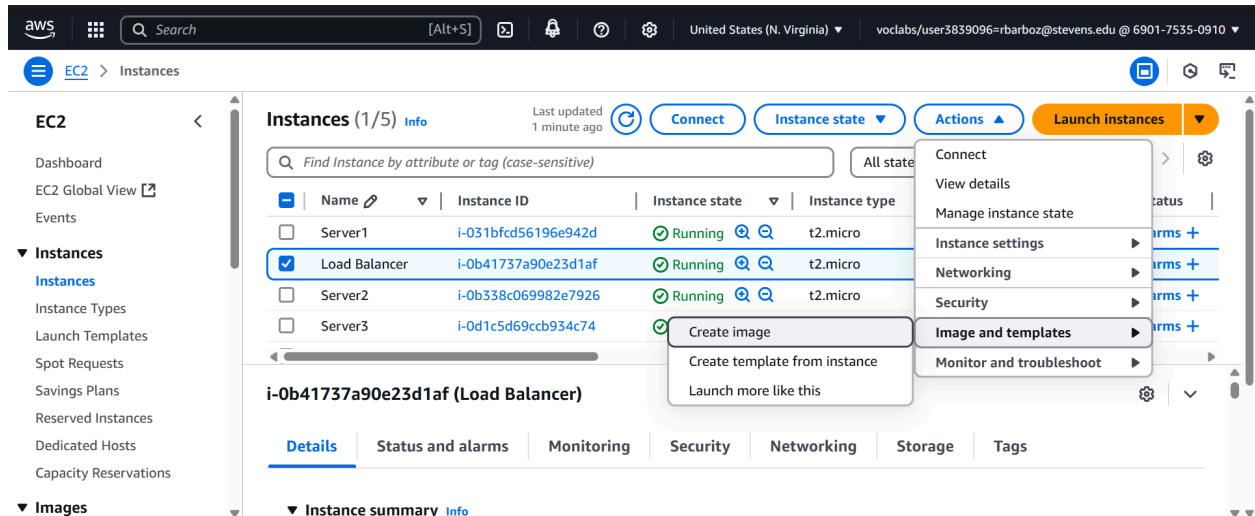
7. To analyze the traffic, I used the tcpdump command (sudo tcpdump -i eth0 -w traffic.pcap) to capture all network packets exchanged between the load balancer and the web servers. This allowed me to observe the types of protocols being used and ensure proper communication between the instances.

```
[ec2-user@ip-172-31-88-116 ~]$ sudo yum install tcpdump -y
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
amzn2-core
Package 14:tcpdump-4.9.2-4.amzn2.1.0.1.x86_64 already installed and latest version
Nothing to do
[ec2-user@ip-172-31-88-116 ~]$ sudo tcpdump -i eth0 -w traffic.pcap
tcpdump: listening on eth0, link-type EN10MB (Ethernet), capture size 262144 bytes
^C86 packets captured
86 packets received by filter
0 packets dropped by kernel
[ec2-user@ip-172-31-88-116 ~]$ sudo tcpdump -r traffic.pcap
reading from file traffic.pcap, link-type EN10MB (Ethernet)
21:18:29.609014 IP ip-172-31-88-116.ec2.internal.ssh > pool-100-35-52-54.nwrknj.fios.verizon.net.62694: Flags [P.], seq 1745134480:1745134524, ac
th 44
21:18:29.609066 IP ip-172-31-88-116.ec2.internal.ssh > pool-100-35-52-54.nwrknj.fios.verizon.net.62694: Flags [P.], seq 44:160, ack 1, win 452, l
21:18:29.627129 IP pool-100-35-52-54.nwrknj.fios.verizon.net.62694 > ip-172-31-88-116.ec2.internal.ssh: Flags [.], ack 160, win 253, length 0
21:18:32.840595 ARP, Request who-has ip-172-31-88-116.ec2.internal tell ip-172-31-80-1.ec2.internal, length 28
21:18:32.840609 ARP, Reply ip-172-31-88-116.ec2.internal is-at 12:10:b5:7c:97:e9 (oui Unknown), length 28
21:18:39.143594 IP ip-172-31-88-116.ec2.internal.50356 > rn-02.koehn.com.ntp: NTPv4, Client, length 48
21:18:39.159883 IP rn-02.koehn.com.ntp > ip-172-31-88-116.ec2.internal.50356: NTPv4, Server, length 48
21:18:40.786224 IP ip-172-31-88-116.ec2.internal.42120 > 169.254.169.123.ntp: NTPv4, Client, length 48
21:18:40.786925 IP 169.254.169.123.ntp > ip-172-31-88-116.ec2.internal.42120: NTPv4, Server, length 48
21:18:56.839582 IP ip-172-31-88-116.ec2.internal.53121 > 169.254.169.123.ntp: NTPv4, Client, length 48
21:18:56.840127 IP 169.254.169.123.ntp > ip-172-31-88-116.ec2.internal.53121: NTPv4, Server, length 48
21:19:02.047522 ARP, Request who-has ip-172-31-80-1.ec2.internal tell ip-172-31-88-116.ec2.internal, length 28
21:19:02.047820 ARP, Reply ip-172-31-80-1.ec2.internal is-at 12:74:3c:d2:04:55 (oui Unknown), length 28
21:19:11.137569 IP pool-100-35-52-54.nwrknj.fios.verizon.net.64009 > ip-172-31-88-116.ec2.internal.http: Flags [S], seq 976386369, win 65535, opt
,nop,nop,sackOK], length 0
21:19:11.137593 IP ip-172-31-88-116.ec2.internal.http > pool-100-35-52-54.nwrknj.fios.verizon.net.64009: Flags [S.], seq 1544814859, ack 97638637
8961,nop,nop,sackOK,nop,wscale 7], length 0
21:19:11.145862 IP pool-100-35-52-54.nwrknj.fios.verizon.net.64010 > ip-172-31-88-116.ec2.internal.http: Flags [S], seq 2708787082, win 65535, op
8,nop,nop,sackOK], length 0
21:19:11.145872 IP ip-172-31-88-116.ec2.internal.http > pool-100-35-52-54.nwrknj.fios.verizon.net.64010: Flags [S.], seq 3242955271, ack 27087870
8961,nop,nop,sackOK,nop,wscale 7], length 0
21:19:11.153600 IP pool-100-35-52-54.nwrknj.fios.verizon.net.64009 > ip-172-31-88-116.ec2.internal.http: Flags [.], ack 1, win 255, length 0
21:19:11.163135 IP pool-100-35-52-54.nwrknj.fios.verizon.net.64010 > ip-172-31-88-116.ec2.internal.http: Flags [.], ack 1, win 255, length 0
21:19:13.030471 IP ip-172-31-88-116.ec2.internal.33732 > 169.254.169.123.ntp: NTPv4, Client, length 48
```

Packet Capture Analysis Summary:

The captured packets show various types of network activity occurring on the EC2 instance. SSH communication is observed between the instance and a remote host, indicating secure remote access. ARP requests and replies show the process of the instance discovering MAC addresses of other devices on the local network. The system is also interacting with NTP servers to synchronize its time. HTTP requests reveal client-server communication over port 80, which is typical for web servers. DHCP traffic shows the instance requesting and receiving an IP address, ensuring it's connected to the network. Additionally, IPv6 DHCP requests indicate that the system is also seeking an IPv6 address. The TCP flags (SYN, ACK, FIN) demonstrate the regular connection setup and teardown processes. Overall, the packet capture confirms that the EC2 instance is performing standard network operations, including secure communication, time synchronization, and network configuration.

8. To create a backup of the load balancer instance, I registered an AMI from the instance. I navigated to the EC2 console, selected the load balancer instance, and created an image of it. This image will serve as a backup and can be used to restore the instance if needed.



9. I launched a new EC2 instance using the AMI I created. This new instance was configured with the same settings and data as the original load balancer instance, ensuring that it can serve as a backup or replacement in case of failure. After launching the new instance, I verified that all files and configurations were intact

aws

Search

[Alt+S]

United States (N. Virginia)

voclabs/user3839096=rbarboz@stevens.edu @ 6901-7535-0910

EC2

Instances

Launch an instance

Launch an instance

Info

Amazon EC2 allows you to create virtual machines, or instances, that run on the AWS Cloud. Quickly get started by following the simple steps below.

Name and tags

Info

Name

royce_lb

Add additional tags

▼ Application and OS Images (Amazon Machine Image)

Info

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Search our full catalog including 1000s of application and OS images

Recents

My AMIs

Quick Start

▼ Summary

Number of instances

Info

1

Software Image (AMI)

royce_load-balancer-backup

ami-0322d5ebde3ba2d45

Virtual server type (instance type)

t2.micro

Firewall (security group)

New security group

Cancel

Launch instance

Preview code

Key pair name - *required*

Royce_lab5

Create new key pair

Network settings

Network

vpc-04b23a6a95db17cc0

Subnet

No preference (Default subnet in any availability zone)

Auto-assign public IP

Enable

Additional charges apply when outside of free tier allowance

Firewall (security groups)

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

Create security group

Select existing security group

We'll create a new security group called 'launch-wizard-4' with the following rules:

Allow SSH traffic from

Helps you connect to your instance

Anywhere

0.0.0.0/0

Allow HTTPS traffic from the internet

To set up an endpoint, for example when creating a web server

Allow HTTP traffic from the internet

To set up an endpoint, for example when creating a web server

Instances (6)

Last updated less than a minute ago

Connect

Instance state

Actions

Launch instances

Find Instance by attribute or tag (case-sensitive)

All states

	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public
<input type="checkbox"/>	Server1	i-031bfcd56196e942d	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1c	ec2-54-85-54-122.com...	54.85.
<input type="checkbox"/>	Load Balancer	i-0b41737a90e23d1af	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1c	ec2-54-82-80-57.comp...	54.82.
<input type="checkbox"/>	Server2	i-0b338c069982e7926	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1c	ec2-13-219-86-50.com...	13.219.
<input type="checkbox"/>	Server3	i-0d1c5d69ccb934c74	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1c	ec2-54-161-55-70.com...	54.161.
<input type="checkbox"/>	Server4	i-0b87449644e3155c8	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1c	ec2-3-95-19-202.comp...	3.95.19.
<input type="checkbox"/>	royce_lb	i-0fa021a810a0460bb	Running	t2.micro	Initializing	View alarms +	us-east-1d	ec2-54-173-232-149.co...	54.173.


```
ec2-user@ip-172-31-19-123:~
GNU nano 2.9.8

events {
    worker_connections 768;
}

http {
    upstream myapp {
        server ec2-54-85-54-122.compute-1.amazonaws.com weight=1;
        server ec2-13-219-86-50.compute-1.amazonaws.com weight=2;
        server ec2-54-161-55-70.compute-1.amazonaws.com weight=1;
        server ec2-3-95-19-202.compute-1.amazonaws.com weight=2;
    }

    server {
        listen 80;
        server_name myapp.com;

        location / {
            proxy_pass http://myapp;
        }
    }
}
```

Terminated all the resources (EC2, security groups, key pairs, custom AMI from snapshot of Loadbalancer):

The screenshot shows the AWS Management Console interface for the EC2 Instances page. The top navigation bar includes the AWS logo, a search bar, and the current region (United States (N. Virginia)). The left sidebar shows the navigation menu with options like Dashboard, EC2 Global View, Events, Instances, Instance Types, Launch Templates, Spot Requests, Savings Plans, Reserved Instances, Dedicated Hosts, Capacity Reservations, Images, AMIs, AMI Catalog, Elastic Block Store, Volumes, Snapshots, and Lifecycle Manager. The main content area displays a table of 6 instances, all in the 'Running' state. The instances are: royce_lb, Server1, Load Balancer, Server2, Server3, and Server4. A dropdown menu is open for the 'Load Balancer' instance, showing options: Stop instance, Start instance, Reboot instance, Hibernate instance, and Terminate (delete) instance. The bottom of the page shows '6 instances selected' and a 'Monitoring' tab.

Name	Instance ID	Instance state	Instance type	Status check	Alarm state	Public IPv4 DNS	Public	
royce_lb	i-0fa021a810a0460bb	Running	t2.micro	2/2 checks passed	View alarms	ec2-54-173-232-149.co...	54.173...	
Server1	i-031bfcd56196e942d	Running	t2.micro	2/2 checks passed	View alarms	us-east-1c	ec2-54-85-54-122.com...	54.85...
Load Balancer	i-0b41737a90e23d1af	Running	t2.micro	2/2 checks passed	View alarms	us-east-1c	ec2-54-82-80-57.comp...	54.82...
Server2	i-0b538c069982e7926	Running	t2.micro	2/2 checks passed	View alarms	us-east-1c	ec2-13-219-86-50.com...	13.219...
Server3	i-0d1c5d69ccb934c74	Running	t2.micro	2/2 checks passed	View alarms	us-east-1c	ec2-54-161-55-70.com...	54.161...
Server4	i-0b87449644e3155c8	Running	t2.micro	2/2 checks passed	View alarms	us-east-1c	ec2-3-95-19-202.comp...	3.95...

Successfully initiated termination (deletion) of i-0fa021a810a0460bb,i-031bfd56196e942d,i-0b41737a90e23d1af,i-0b338c069982e7926,i-0d1c5d69ccb934c74,i-0b87449644e3155c8

Instances (6) Info

Last updated less than a minute ago

Connect

Instance state

Actions

Launch instances

Find Instance by attribute or tag (case-sensitive)

All states

< 1 >

<input type="checkbox"/>	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public
<input type="checkbox"/>	Server1	i-031bfd56196e942d	Terminated	t2.micro	-	View alarms +	us-east-1c	-	-
<input type="checkbox"/>	Load Balancer	i-0b41737a90e23d1af	Terminated	t2.micro	-	View alarms +	us-east-1c	-	-
<input type="checkbox"/>	Server2	i-0b338c069982e7926	Terminated	t2.micro	-	View alarms +	us-east-1c	-	-
<input type="checkbox"/>	Server3	i-0d1c5d69ccb934c74	Terminated	t2.micro	-	View alarms +	us-east-1c	-	-
<input type="checkbox"/>	Server4	i-0b87449644e3155c8	Terminated	t2.micro	-	View alarms +	us-east-1c	-	-
<input type="checkbox"/>	royce_lb	i-0fa021a810a0460bb	Terminated	t2.micro	-	View alarms +	us-east-1d	-	-

Successfully deleted 4 security groups

Details

Security Groups (1) Info

Actions

Export security groups to CSV

Create security group

Find resources by attribute or tag

< 1 >

<input type="checkbox"/>	Name	Security group ID	Security group name	VPC ID	Description
<input type="checkbox"/>	-	sg-069bd052e521befbb	default	vpc-04b23a6a95db17cc0	default VPC security group

Successfully deleted 2 key pairs

Notifications

Key pairs Info

Actions

Create key pair

Find Key Pair by attribute or tag

< 1 >

<input checked="" type="checkbox"/>	Name	Type	Created	Fingerprint	ID
-------------------------------------	------	------	---------	-------------	----

No key pairs to display

Amazon Machine Images (AMIs) (1/1) Info

Recycle Bin

EC2 Image Builder

Actions

Launch instance from AMI

Owned by me

Find AMI by attribute or tag

< 1 >

<input checked="" type="checkbox"/>	Name	AMI name	AMI ID	Source	Owner
<input checked="" type="checkbox"/>	royce_load-balancer-backup	ami-0322d5ebde3ba2d45	690175350910/royce_load-balancer-ba...	690175350910	

AMI ID: ami-0322d5ebde3ba2d45

DetailsPermissionsStorageTags

AMI ID

ami-0322d5ebde3ba2d45

Image type

machine

Platform details

Linux/UNIX

Root device type

EBS

AMI name

royce_load-balancer-backup

Owner account ID

690175350910

Architecture

x86_64

Usage operation

RunInstances

Root device name

/dev/xvda

Status

Available

Source

690175350910/royce_load-balancer-backup

Virtualization type

hvm

Successfully deregistered ami-0322d5ebde3ba2d45.

Amazon Machine Images (AMIs)

Info

Owned by me

Find AMI by attribute or tag

Recycle Bin

EC2 Image Builder

Actions

Launch instance from AMI

Owned by me

Find AMI by attribute or tag

Name

AMI name

AMI ID

Source

Owner

Visibility

Status

No AMIs in this Region for: Owned by me.
You can use the filter to view Owned By Me, Private Images, Public Images, or Disabled Images.

Successfully deleted snapshot snap-00ff1e7c7bd5ad17e.

Snapshots

Info

Owned by me

Search

Recycle Bin

Actions

Create snapshot

Owned by me

Search

Name

Snapshot ID

Full snapshot size

Volume size

Description

Storage tier

Snapshot status

Started

You currently have no snapshots in this Region.

aws

Search

United States (N. Virginia)

vociabo/user3839096=rbarboz@stevens.edu @ 6901-7535-0910

EC2

Dashboard

EC2 Global View

Events

Instances

Instance Types

Launch Templates

Spot Requests

Resources

You are using the following Amazon EC2 resources in the United States (N. Virginia) Region:

Instances (running)

0

Auto Scaling Groups

0

Capacity Reservations

0

Dedicated Hosts

0

Elastic IPs

0

Instances

6

Key pairs

0

Load balancers

0

Placement groups

0

Security groups

1

Snapshots

0

Volumes

0

Account attributes

Default VPC

vpc-04b23a6a95db17cc0

Settings

Data protection and security

Allowed AMIs

Zones

EC2 Serial Console

Default credit specification